

ytdtz6oz6

November 13, 2025

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
```

1 LendSmart Credit Risk Analysis

Paulina Leal Mosqueda, A01659576

Santiago Nava Figueroa, A01174557

2 Section 1

```
[ ]: # Importar csv

df = pd.read_csv('credit_risk_data-1.csv')
df.head()
```



```
[ ]:   application_id application_date  loan_amount  annual_income \
0          APP_2328      2022-01-01    132221.82     60451.82
1          APP_558       2022-01-01    134906.42     114634.08
2          APP_2477      2022-01-01     30285.19     82772.53
3          APP_741       2022-01-01    32516.09     94023.36
4          APP_145       2022-01-02    77900.99     53515.02

      employment_years  job_stability_score  credit_score  credit_utilization \
0                  6.6           0.898         679           0.106
1                 10.3           0.808         718           0.030
2                 12.1           0.964         768           0.174
3                  9.1           0.690         670           0.141
4                  7.2           0.679         651           0.097

      payment_history_score  open_credit_lines  debt_to_income_ratio \
0                  0.876                   1                0.451
```

| | | | | | | |
|---|-----------------------|-------------|-------|-----------------|----------------|---|
| 1 | 0.719 | 4 | 0.090 | | | |
| 2 | 0.775 | 6 | 0.201 | | | |
| 3 | 0.993 | 3 | 0.322 | | | |
| 4 | 0.946 | 2 | 0.222 | | | |
| | savings_ratio | asset_value | age | education_level | marital_status | \ |
| 0 | 0.500 | 352569.55 | 41 | High School | Married | |
| 1 | 0.235 | 224364.21 | 46 | Masters | Divorced | |
| 2 | 0.172 | 514765.55 | 44 | High School | Widowed | |
| 3 | 0.368 | 182541.72 | 26 | Bachelors | Single | |
| 4 | 0.324 | 223691.29 | 50 | Associates | Single | |
| | residential_stability | loan_status | | | | |
| 0 | 3.5 | 0 | | | | |
| 1 | 11.4 | 0 | | | | |
| 2 | 8.6 | 0 | | | | |
| 3 | 3.9 | 0 | | | | |
| 4 | 9.6 | 0 | | | | |

[]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   application_id    2500 non-null   object 
 1   application_date  2500 non-null   object 
 2   loan_amount       2500 non-null   float64
 3   annual_income     2500 non-null   float64
 4   employment_years  2500 non-null   float64
 5   job_stability_score  2500 non-null   float64
 6   credit_score      2500 non-null   int64  
 7   credit_utilization 2500 non-null   float64
 8   payment_history_score  2500 non-null   float64
 9   open_credit_lines  2500 non-null   int64  
 10  debt_to_income_ratio  2500 non-null   float64
 11  savings_ratio     2500 non-null   float64
 12  asset_value       2500 non-null   float64
 13  age               2500 non-null   int64  
 14  education_level   2500 non-null   object 
 15  marital_status    2500 non-null   object 
 16  residential_stability  2500 non-null   float64
 17  loan_status       2500 non-null   int64  
dtypes: float64(10), int64(4), object(4)
memory usage: 351.7+ KB
```

```
[ ]: df.describe()

[ ]:      loan_amount  annual_income  employment_years  job_stability_score \
count    2500.000000    2500.000000    2500.000000    2500.000000
mean    155716.305344   67707.807596     6.675640    0.634643
std     149605.357952   27302.931731     3.488021    0.293276
min     5000.000000   15000.000000     0.000000    0.011000
25%    42984.517500   47475.317500     4.000000    0.375500
50%    97054.315000   66963.475000     6.700000    0.752000
75%   213214.992500   87347.642500     9.300000    0.866000
max    500000.000000  149929.960000    19.300000    0.999000

      credit_score  credit_utilization  payment_history_score \
count    2500.000000    2500.000000    2500.000000
mean    681.728400     0.358176     0.740733
std     88.683309     0.289995     0.285966
min     334.000000     0.004000     0.029000
25%    642.750000     0.131000     0.517500
50%    700.000000     0.246000     0.880500
75%    743.000000     0.592250     0.956000
max    850.000000     0.998000     1.000000

      open_credit_lines  debt_to_income_ratio  savings_ratio  asset_value \
count    2500.000000    2500.000000    2500.000000    2500.000000
mean     3.451600     0.408094     0.320784    175666.741236
std     2.083793     0.224736     0.192079    182652.568930
min     0.000000     0.009000     0.000000    550.630000
25%    2.000000     0.228000     0.161000    49513.082500
50%    3.000000     0.359000     0.327000   121018.750000
75%    5.000000     0.565000     0.464000   235513.902500
max    11.000000     0.979000     0.893000   1000000.000000

      age  residential_stability  loan_status
count  2500.000000    2500.000000    2500.000000
mean   42.045600     6.023200     0.265600
std    12.092395     3.205397     0.441741
min    18.000000     0.000000     0.000000
25%   34.000000     3.600000     0.000000
50%   42.000000     5.900000     0.000000
75%   50.000000     8.400000     1.000000
max   75.000000    16.400000     1.000000

[ ]: # find nan values
df.isnull().sum()

[ ]: application_id          0
application_date           0
```

```
loan_amount          0
annual_income        0
employment_years     0
job_stability_score 0
credit_score         0
credit_utilization   0
payment_history_score 0
open_credit_lines    0
debt_to_income_ratio 0
savings_ratio        0
asset_value          0
age                  0
education_level      0
marital_status        0
residential_stability 0
loan_status           0
dtype: int64
```


3 Section 2

```
[ ]: # Plot the distribution of loan_status and calculate the exact default rate.

loan_status = df['loan_status']

plt.figure(figsize=(8, 6))
# histogram
sns.countplot(x=loan_status)
plt.title('Distribution of Loan Status')
plt.xlabel('Loan Status (0: No Default, 1: Default)')
plt.ylabel('Count')
plt.show()

# Calculate the exact default rate.
total_loans = len(loan_status)
default_count = loan_status.value_counts()[1]
default_rate = (default_count / total_loans) * 100
print(f"Total Loans: {total_loans}")
print(f"Default Count: {default_count}")
print(f"Default Rate: {default_rate:.2f}%)
```



Total Loans: 2500
Default Count: 664
Default Rate: 26.56%

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Continuous variables to visualize
continuous_vars = [
    'credit_score', 'annual_income', 'debt_to_income_ratio', 'loan_amount', 'employment_years',
    'job_stability_score', 'credit_utilization', 'payment_history_score',
    'savings_ratio', 'asset_value', 'age', 'residential_stability'
]

# Set plot style
sns.set(style="whitegrid")

plt.figure(figsize=(15, 20))
```

```

for i, var in enumerate(continuous_vars, 1):
    plt.subplot(6, 2, i)
    sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')
    plt.title(f'{var} by Loan Status')
    plt.xlabel('Loan Status (0 = Non-Default, 1 = Default)')
    plt.ylabel(var)

plt.tight_layout()
plt.show()

plt.tight_layout()
plt.show()

```

/tmp/ipython-input-603495265.py:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')

```

/tmp/ipython-input-603495265.py:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')

```

/tmp/ipython-input-603495265.py:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')

```

/tmp/ipython-input-603495265.py:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')

```

/tmp/ipython-input-603495265.py:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')
/tmp/ipython-input-603495265.py:18: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in
v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same
effect.

sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')
/tmp/ipython-input-603495265.py:18: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in
v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same
effect.


sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')
/tmp/ipython-input-603495265.py:18: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in
v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same
effect.

sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')
/tmp/ipython-input-603495265.py:18: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in
v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same
effect.

sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')
/tmp/ipython-input-603495265.py:18: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in
v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same
effect.

sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')
/tmp/ipython-input-603495265.py:18: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in
v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same
effect.
```

```
sns.boxplot(data=df, x='loan_status', y=var, palette='Set2')
```




```
<Figure size 640x480 with 0 Axes>
```

```
[ ]: # Gráficar continous variables com histograms
```

```
plt.figure(figsize=(15, 20))
for i, var in enumerate(continuous_vars, 1):
    plt.subplot(6, 2, i)
    sns.histplot(data=df, x=var, kde=True, color='skyblue')
    plt.title(f'Distribución de {var}')
    plt.xlabel(var)
    plt.ylabel('Frequency')

plt
plt.tight_layout()
plt.show()
```



```
[ ]: df['education_level'].value_counts()
```

```
# groupby educational level for mean default rate
df.groupby('education_level')['loan_status'].mean()
```

```
[ ]: education_level
Associates      0.369128
Bachelors       0.154676
Doctorate        0.048913
High School     0.560811
Masters          0.128959
Name: loan_status, dtype: float64
```

```
[ ]: # Set style
sns.set(style="whitegrid")

default_rates_educ = (
    df.groupby('education_level')['loan_status']
    .mean()
    .reset_index()
    .rename(columns={'loan_status': 'default_rate'})
)



default_rates_marital = (
    df.groupby('marital_status')['loan_status']
    .mean()
    .reset_index()
    .rename(columns={'loan_status': 'default_rate'})
)
```



```
plt.figure(figsize=(10, 5))
sns.barplot(data=default_rates_educ, x='education_level', y='default_rate',
            hue='education_level', legend=False, palette='viridis')
plt.title('Mean Default Rate by Education Level')
plt.xlabel('Education Level')
plt.ylabel('Mean Default Rate')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

plt.figure(figsize=(8, 5))
sns.barplot(data=default_rates_marital, x='marital_status', y='default_rate',
            hue='marital_status', legend=False, palette='magma')
plt.title('Mean Default Rate by Marital Status')
plt.xlabel('Marital Status')
plt.ylabel('Mean Default Rate')
```

```
plt.tight_layout()  
plt.show()
```




```
[ ]: # Correltaion for continuous_vars  
  
correlation_matrix = df[continuous_vars].corr()
```

```

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Matrix')
plt.show()

```



```
[ ]: correlation_matrix.describe()
```

```

[ ]:   credit_score  annual_income  debt_to_income_ratio  loan_amount \
count    12.000000      12.000000          12.000000      12.000000
mean     0.280763      0.233603         -0.232651      0.079472
std      0.487464      0.389996          0.548401      0.290230
min     -0.657600     -0.467363         -0.732091     -0.034981
25%      0.171627      0.138400         -0.551312     -0.012127
50%      0.379584      0.298325         -0.403739     -0.000569
75%      0.550242      0.390495         -0.227114      0.011519
max      1.000000      1.000000          1.000000      1.000000

```

| | | | | |
|-------|-----------------------|---------------------|--------------------|-----------|
| | employment_years | job_stability_score | credit_utilization | \ |
| count | 12.000000 | 12.000000 | 12.000000 | |
| mean | 0.258679 | 0.319335 | -0.287300 | |
| std | 0.443269 | 0.568965 | 0.584199 | |
| min | -0.563996 | -0.817940 | -0.845028 | |
| 25% | 0.165648 | 0.237627 | -0.634631 | |
| 50% | 0.324901 | 0.475626 | -0.476804 | |
| 75% | 0.488187 | 0.643857 | -0.243293 | |
| max | 1.000000 | 1.000000 | 1.000000 | |
| | payment_history_score | savings_ratio | asset_value | age \ |
| count | 12.000000 | 12.000000 | 12.000000 | 12.000000 |
| mean | 0.323260 | 0.279506 | 0.214222 | 0.190159 |
| std | 0.578606 | 0.472911 | 0.358205 | 0.338189 |
| min | -0.845028 | -0.626975 | -0.380645 | -0.328598 |
| 25% | 0.251532 | 0.195030 | 0.121195 | 0.117523 |
| 50% | 0.478424 | 0.358970 | 0.249944 | 0.213832 |
| 75% | 0.663408 | 0.542854 | 0.329894 | 0.272843 |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| | residential_stability | | | |
| count | 12.000000 | | | |
| mean | 0.237749 | | | |
| std | 0.400431 | | | |
| min | -0.486246 | | | |
| 25% | 0.150872 | | | |
| 50% | 0.293684 | | | |
| 75% | 0.416695 | | | |
| max | 1.000000 | | | |

4 Variable con alta multiconlinealidad siendo 0.7

Credit utilization y job stability score - 0.82

job stability score y debt to income ratio - 0.72

Credit utilization y payment history score - 0.85

debt to income ratio y Credit utilization - 0.71

payment history score y job stability score - 0.85

debt to income ratio y payment history score - 0.73

Sin embargo para elegir las variables se usara RFE para tomar las variables con baja multicolinealidad.

5 Section 3

```
[ ]: df.education_level.unique()
```

```
[ ]: array(['High School', 'Masters', 'Bachelors', 'Associates', 'Doctorate'],
      dtype=object)
```

```
[ ]: df.marital_status.unique()
```

```
[ ]: array(['Married', 'Divorced', 'Widowed', 'Single'], dtype=object)
```

```
[ ]: educationMAP = {'High School':1, 'Masters':4, 'Bachelors':3, 'Associates':2, 'Doctorate':5}
df['education_level'] = df['education_level'].map(educationMAP)
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   application_id    2500 non-null   object 
 1   application_date  2500 non-null   object 
 2   loan_amount       2500 non-null   float64
 3   annual_income     2500 non-null   float64
 4   employment_years  2500 non-null   float64
 5   job_stability_score 2500 non-null   float64
 6   credit_score      2500 non-null   int64  
 7   credit_utilization 2500 non-null   float64
 8   payment_history_score 2500 non-null   float64
 9   open_credit_lines  2500 non-null   int64  
 10  debt_to_income_ratio 2500 non-null   float64
 11  savings_ratio     2500 non-null   float64
 12  asset_value       2500 non-null   float64
 13  age                2500 non-null   int64  
 14  education_level   2500 non-null   int64  
 15  marital_status    2500 non-null   object 
 16  residential_stability 2500 non-null   float64
 17  loan_status       2500 non-null   int64  
dtypes: float64(10), int64(5), object(3)
memory usage: 351.7+ KB
```

```
[ ]: df_encoded = pd.get_dummies(df, columns=['marital_status'], dtype = float)
```

```
df_encoded = df_encoded[['application_id', 'application_date', 'loan_amount', 'annual_income',
```

```

'employment_years', 'job_stability_score', 'credit_score',
'credit_utilization', 'payment_history_score', 'open_credit_lines',
'debt_to_income_ratio', 'savings_ratio', 'asset_value', 'age',
'education_level', 'residential_stability',
'marital_status_Divorced', 'marital_status_Married',
'marital_status_Single', 'marital_status_Widowed', "loan_status"]]

df_encoded.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   application_id    2500 non-null   object  
 1   application_date  2500 non-null   object  
 2   loan_amount       2500 non-null   float64 
 3   annual_income     2500 non-null   float64 
 4   employment_years  2500 non-null   float64 
 5   job_stability_score  2500 non-null   float64 
 6   credit_score      2500 non-null   int64   
 7   credit_utilization 2500 non-null   float64 
 8   payment_history_score  2500 non-null   float64 
 9   open_credit_lines  2500 non-null   int64   
 10  debt_to_income_ratio  2500 non-null   float64 
 11  savings_ratio     2500 non-null   float64 
 12  asset_value       2500 non-null   float64 
 13  age               2500 non-null   int64   
 14  education_level   2500 non-null   int64   
 15  residential_stability  2500 non-null   float64 
 16  marital_status_Divorced  2500 non-null   float64 
 17  marital_status_Married  2500 non-null   float64 
 18  marital_status_Single  2500 non-null   float64 
 19  marital_status_Widowed  2500 non-null   float64 
 20  loan_status        2500 non-null   int64   

dtypes: float64(14), int64(5), object(2)
memory usage: 410.3+ KB

```

```
[ ]: X = df_encoded.iloc[:,3:20]
```

```

X.info()
y = df_encoded.iloc[:,-1]
y

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   application_id    2500 non-null   object  
 1   application_date  2500 non-null   object  
 2   loan_amount       2500 non-null   float64 
 3   annual_income     2500 non-null   float64 
 4   employment_years  2500 non-null   float64 
 5   job_stability_score  2500 non-null   float64 
 6   credit_score      2500 non-null   int64   
 7   credit_utilization 2500 non-null   float64 
 8   payment_history_score  2500 non-null   float64 
 9   open_credit_lines  2500 non-null   int64   
 10  debt_to_income_ratio  2500 non-null   float64 
 11  savings_ratio     2500 non-null   float64 
 12  asset_value       2500 non-null   float64 
 13  age               2500 non-null   int64   
 14  education_level   2500 non-null   int64   
 15  residential_stability  2500 non-null   float64 
 16  marital_status_Divorced  2500 non-null   float64 
 17  marital_status_Married  2500 non-null   float64 
 18  marital_status_Single  2500 non-null   float64 
 19  marital_status_Widowed  2500 non-null   float64 
 20  loan_status        2500 non-null   int64   

dtypes: float64(14), int64(5), object(2)
memory usage: 410.3+ KB

```

```
0    annual_income            2500 non-null   float64
1    employment_years         2500 non-null   float64
2    job_stability_score     2500 non-null   float64
3    credit_score             2500 non-null   int64
4    credit_utilization      2500 non-null   float64
5    payment_history_score   2500 non-null   float64
6    open_credit_lines        2500 non-null   int64
7    debt_to_income_ratio    2500 non-null   float64
8    savings_ratio            2500 non-null   float64
9    asset_value              2500 non-null   float64
10   age                      2500 non-null   int64
11   education_level          2500 non-null   int64
12   residential_stability   2500 non-null   float64
13   marital_status_Divorced 2500 non-null   float64
14   marital_status_Married   2500 non-null   float64
15   marital_status_Single    2500 non-null   float64
16   marital_status_Widowed   2500 non-null   float64
dtypes: float64(13), int64(4)
memory usage: 332.2 KB
```

```
[ ]: 0      0
1      0
2      0
3      0
4      0
..
2495  0
2496  1
2497  0
2498  1
2499  0
Name: loan_status, Length: 2500, dtype: int64
```

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,           # 20% para prueba
    random_state=42,         # Semilla para reproducibilidad
    stratify=y               # Opcional: mantiene proporciones de clases (útil en
                             # clasificación)
)
```

```
[ ]: X.columns
```

```
[ ]: Index(['annual_income', 'employment_years', 'job_stability_score',
       'credit_score', 'credit_utilization', 'payment_history_score',
       'open_credit_lines', 'debt_to_income_ratio', 'savings_ratio',
       'asset_value', 'age', 'education_level', 'residential_stability',
```

```
'marital_status_Divorced', 'marital_status_Married',
'marital_status_Single', 'marital_status_Widowed'],
dtype='object')
```

```
[ ]: scaler = MinMaxScaler(feature_range=(0, 1))
columnas_a_escalar =['annual_income', 'employment_years',
    'credit_score',
    'open_credit_lines', 'debt_to_income_ratio',
    'savings_ratio', 'asset_value', 'age',,
    ↴'education_level','residential_stability']
```

```
X_scaled_train = pd.DataFrame(
    scaler.fit_transform(X_train[columnas_a_escalar]),
    columns=columnas_a_escalar,
    index=X_train.index
)
```

```
[ ]: X_train
```

```
[ ]: annual_income employment_years job_stability_score credit_score \
2082      62648.24          9.9        0.541       714
1319     113346.13          8.5        0.929       719
1569      57588.47         13.6        0.769       742
1257      83159.53         12.1        0.799       779
2096      64288.01          6.0        0.519       722
...
1165      45946.99          6.1        0.829       626
1334      86267.61          9.0        0.830       706
626       34526.29          4.5        0.181       545
1901      61908.46          5.0        0.818       643
1373      89359.66          6.5        0.824       695

    credit_utilization payment_history_score open_credit_lines \
2082           0.171            0.815             3
1319           0.219            0.938             4
1569           0.244            0.815             1
1257           0.046            0.732             4
2096           0.148            0.949             2
...
1165           0.085            0.932             5
1334           0.063            0.994             3
626            0.922            0.261             2
1901           0.222            0.905             4
1373           0.271            0.941             5
```

| | debt_to_income_ratio | savings_ratio | asset_value | age | education_level | \ |
|------|-----------------------|-------------------------|------------------------|-----|-----------------|-----|
| 2082 | 0.240 | 0.549 | 168748.10 | 39 | | 2 |
| 1319 | 0.303 | 0.516 | 272892.03 | 59 | | 4 |
| 1569 | 0.659 | 0.550 | 82093.48 | 41 | | 3 |
| 1257 | 0.646 | 0.402 | 74527.88 | 71 | | 2 |
| 2096 | 0.233 | 0.394 | 122808.69 | 52 | | 5 |
| ... | ... | ... | ... | ... | ... | |
| 1165 | 0.332 | 0.370 | 143994.83 | 26 | | 1 |
| 1334 | 0.462 | 0.392 | 218923.92 | 68 | | 4 |
| 626 | 0.786 | 0.246 | 25799.60 | 45 | | 3 |
| 1901 | 0.234 | 0.262 | 67014.03 | 52 | | 4 |
| 1373 | 0.222 | 0.550 | 97451.56 | 64 | | 5 |
| | | | | | | |
| | residential_stability | marital_status_Divorced | marital_status_Married | | | \ |
| 2082 | 3.1 | | 0.0 | | | 1.0 |
| 1319 | 5.5 | | 0.0 | | | 1.0 |
| 1569 | 3.8 | | 0.0 | | | 1.0 |
| 1257 | 5.8 | | 0.0 | | | 0.0 |
| 2096 | 7.2 | | 1.0 | | | 0.0 |
| ... | ... | ... | ... | ... | ... | |
| 1165 | 7.4 | | 1.0 | | | 0.0 |
| 1334 | 8.1 | | 0.0 | | | 0.0 |
| 626 | 2.9 | | 1.0 | | | 0.0 |
| 1901 | 7.7 | | 0.0 | | | 1.0 |
| 1373 | 10.0 | | 0.0 | | | 0.0 |
| | | | | | | |
| | marital_status_Single | marital_status_Widowed | | | | |
| 2082 | 0.0 | | 0.0 | | | |
| 1319 | 0.0 | | 0.0 | | | |
| 1569 | 0.0 | | 0.0 | | | |
| 1257 | 0.0 | | 1.0 | | | |
| 2096 | 0.0 | | 0.0 | | | |
| ... | ... | ... | ... | ... | ... | |
| 1165 | 0.0 | | 0.0 | | | |
| 1334 | 1.0 | | 0.0 | | | |
| 626 | 0.0 | | 0.0 | | | |
| 1901 | 0.0 | | 0.0 | | | |
| 1373 | 1.0 | | 0.0 | | | |

[2000 rows x 17 columns]

```
[ ]: X_scaled_train.columns = columnas_a_escalar
X_scaled_train
```

```
[ ]: annual_income employment_years credit_score open_credit_lines \
2082      0.353133        0.585799      0.736434      0.272727
1319      0.728868        0.502959      0.746124      0.363636
```

| | | | | |
|------|----------------------|-----------------------|-------------|----------|
| 1569 | 0.315634 | 0.804734 | 0.790698 | 0.090909 |
| 1257 | 0.505147 | 0.715976 | 0.862403 | 0.363636 |
| 2096 | 0.365286 | 0.355030 | 0.751938 | 0.181818 |
| ... | ... | ... | ... | ... |
| 1165 | 0.229356 | 0.360947 | 0.565891 | 0.454545 |
| 1334 | 0.528182 | 0.532544 | 0.720930 | 0.272727 |
| 626 | 0.144714 | 0.266272 | 0.408915 | 0.181818 |
| 1901 | 0.347650 | 0.295858 | 0.598837 | 0.363636 |
| 1373 | 0.551098 | 0.384615 | 0.699612 | 0.454545 |
| | debt_to_income_ratio | savings_ratio | asset_value | age \ |
| 2082 | 0.229406 | 0.614782 | 0.168290 | 0.368421 |
| 1319 | 0.295099 | 0.577828 | 0.272491 | 0.719298 |
| 1569 | 0.666319 | 0.615901 | 0.081588 | 0.403509 |
| 1257 | 0.652763 | 0.450168 | 0.074018 | 0.929825 |
| 2096 | 0.222106 | 0.441209 | 0.122325 | 0.596491 |
| ... | ... | ... | ... | ... |
| 1165 | 0.325339 | 0.414334 | 0.143523 | 0.140351 |
| 1334 | 0.460897 | 0.438970 | 0.218494 | 0.877193 |
| 626 | 0.798749 | 0.275476 | 0.025263 | 0.473684 |
| 1901 | 0.223149 | 0.293393 | 0.066500 | 0.596491 |
| 1373 | 0.210636 | 0.615901 | 0.096954 | 0.807018 |
| | education_level | residential_stability | | |
| 2082 | 0.25 | 0.189024 | | |
| 1319 | 0.75 | 0.335366 | | |
| 1569 | 0.50 | 0.231707 | | |
| 1257 | 0.25 | 0.353659 | | |
| 2096 | 1.00 | 0.439024 | | |
| ... | ... | ... | | |
| 1165 | 0.00 | 0.451220 | | |
| 1334 | 0.75 | 0.493902 | | |
| 626 | 0.50 | 0.176829 | | |
| 1901 | 0.75 | 0.469512 | | |
| 1373 | 1.00 | 0.609756 | | |

[2000 rows x 10 columns]

```
[ ]: X_train = X_train.drop(columns = columnas_a_escalar)
X_train
```

```
[ ]: job_stability_score credit_utilization payment_history_score \
2082          0.541           0.171          0.815
1319          0.929           0.219          0.938
1569          0.769           0.244          0.815
1257          0.799           0.046          0.732
2096          0.519           0.148          0.949
```

| | | | |
|--|-------|-------|-------|
| ... | ... | ... | ... |
| 1165 | 0.829 | 0.085 | 0.932 |
| 1334 | 0.830 | 0.063 | 0.994 |
| 626 | 0.181 | 0.922 | 0.261 |
| 1901 | 0.818 | 0.222 | 0.905 |
| 1373 | 0.824 | 0.271 | 0.941 |
| marital_status_Divorced marital_status_Married marital_status_Single \ | | | |
| 2082 | 0.0 | 1.0 | 0.0 |
| 1319 | 0.0 | 1.0 | 0.0 |
| 1569 | 0.0 | 1.0 | 0.0 |
| 1257 | 0.0 | 0.0 | 0.0 |
| 2096 | 1.0 | 0.0 | 0.0 |
| ... | ... | ... | ... |
| 1165 | 1.0 | 0.0 | 0.0 |
| 1334 | 0.0 | 0.0 | 1.0 |
| 626 | 1.0 | 0.0 | 0.0 |
| 1901 | 0.0 | 1.0 | 0.0 |
| 1373 | 0.0 | 0.0 | 1.0 |
| marital_status_Widowed | | | |
| 2082 | 0.0 | | |
| 1319 | 0.0 | | |
| 1569 | 0.0 | | |
| 1257 | 1.0 | | |
| 2096 | 0.0 | | |
| ... | ... | | |
| 1165 | 0.0 | | |
| 1334 | 0.0 | | |
| 626 | 0.0 | | |
| 1901 | 0.0 | | |
| 1373 | 0.0 | | |

[2000 rows x 7 columns]

```
[ ]: #X_scaled_train = X_scaled_train.reset_index(drop=True)
#X_train = X_train.reset_index(drop=True)

X_train = pd.concat([X_scaled_train, X_train], axis=1)
X_train
```

```
[ ]: annual_income employment_years credit_score open_credit_lines \
2082      0.353133      0.585799      0.736434      0.272727
1319      0.728868      0.502959      0.746124      0.363636
1569      0.315634      0.804734      0.790698      0.090909
1257      0.505147      0.715976      0.862403      0.363636
2096      0.365286      0.355030      0.751938      0.181818
```

| | | | | |
|------|------------------------|-----------------------|-------------------------|----------|
| ... | ... | ... | ... | ... |
| 1165 | 0.229356 | 0.360947 | 0.565891 | 0.454545 |
| 1334 | 0.528182 | 0.532544 | 0.720930 | 0.272727 |
| 626 | 0.144714 | 0.266272 | 0.408915 | 0.181818 |
| 1901 | 0.347650 | 0.295858 | 0.598837 | 0.363636 |
| 1373 | 0.551098 | 0.384615 | 0.699612 | 0.454545 |
| | | | | |
| | debt_to_income_ratio | savings_ratio | asset_value | age \ |
| 2082 | 0.229406 | 0.614782 | 0.168290 | 0.368421 |
| 1319 | 0.295099 | 0.577828 | 0.272491 | 0.719298 |
| 1569 | 0.666319 | 0.615901 | 0.081588 | 0.403509 |
| 1257 | 0.652763 | 0.450168 | 0.074018 | 0.929825 |
| 2096 | 0.222106 | 0.441209 | 0.122325 | 0.596491 |
| ... | ... | ... | ... | ... |
| 1165 | 0.325339 | 0.414334 | 0.143523 | 0.140351 |
| 1334 | 0.460897 | 0.438970 | 0.218494 | 0.877193 |
| 626 | 0.798749 | 0.275476 | 0.025263 | 0.473684 |
| 1901 | 0.223149 | 0.293393 | 0.066500 | 0.596491 |
| 1373 | 0.210636 | 0.615901 | 0.096954 | 0.807018 |
| | | | | |
| | education_level | residential_stability | job_stability_score | \ |
| 2082 | 0.25 | 0.189024 | | 0.541 |
| 1319 | 0.75 | 0.335366 | | 0.929 |
| 1569 | 0.50 | 0.231707 | | 0.769 |
| 1257 | 0.25 | 0.353659 | | 0.799 |
| 2096 | 1.00 | 0.439024 | | 0.519 |
| ... | ... | ... | ... | ... |
| 1165 | 0.00 | 0.451220 | | 0.829 |
| 1334 | 0.75 | 0.493902 | | 0.830 |
| 626 | 0.50 | 0.176829 | | 0.181 |
| 1901 | 0.75 | 0.469512 | | 0.818 |
| 1373 | 1.00 | 0.609756 | | 0.824 |
| | | | | |
| | credit_utilization | payment_history_score | marital_status_Divorced | \ |
| 2082 | 0.171 | 0.815 | | 0.0 |
| 1319 | 0.219 | 0.938 | | 0.0 |
| 1569 | 0.244 | 0.815 | | 0.0 |
| 1257 | 0.046 | 0.732 | | 0.0 |
| 2096 | 0.148 | 0.949 | | 1.0 |
| ... | ... | ... | ... | ... |
| 1165 | 0.085 | 0.932 | | 1.0 |
| 1334 | 0.063 | 0.994 | | 0.0 |
| 626 | 0.922 | 0.261 | | 1.0 |
| 1901 | 0.222 | 0.905 | | 0.0 |
| 1373 | 0.271 | 0.941 | | 0.0 |
| | | | | |
| | marital_status_Married | marital_status_Single | marital_status_Widowed | |

| | | | |
|------|-----|-----|-----|
| 2082 | 1.0 | 0.0 | 0.0 |
| 1319 | 1.0 | 0.0 | 0.0 |
| 1569 | 1.0 | 0.0 | 0.0 |
| 1257 | 0.0 | 0.0 | 1.0 |
| 2096 | 0.0 | 0.0 | 0.0 |
| ... | ... | ... | ... |
| 1165 | 0.0 | 0.0 | 0.0 |
| 1334 | 0.0 | 1.0 | 0.0 |
| 626 | 0.0 | 0.0 | 0.0 |
| 1901 | 1.0 | 0.0 | 0.0 |
| 1373 | 0.0 | 1.0 | 0.0 |

[2000 rows x 17 columns]

6 Section 4

key statistical assumptions that differentiate LDA and QDA

Antes de aplicar los modelos, se debe determinar las diferencias estadísticas de LDA y QDA.

LDA es un modelo de clasificación que se basa en identificar grupos al considerar que todas las variables tienen una covarianza lineal. Esto quiere decir que se asume que todas las variables provienen de una distribución Gaussiana. Por otra parte, QDA asume que las covarianzas son cuadráticas.

Multivariate Normality: observando los histogramas hechos anteriormente, se puede observar como es que las variables no tienen una distribución Gaussiana en su mayoría.

Homogeneity of Covariance Matrices: utilizando las gráficas anteriores se observa como es que no hay una homogeneidad en las matrices de covarianza, ya que la distribución de cada variable es diferente.

Por lo tanto, se plantea la siguiente hipótesis: Si las matrices de covarianza son diferentes, entonces, se espera que el modelo QDA tenga mejores resultados que LDA.

7 Section 5 - Linear Discriminant Analysis (LDA)

```
[ ]: from sklearn.preprocessing import StandardScaler
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
import pandas as pd
import numpy as np

# Fit LDA model
lda = LinearDiscriminantAnalysis()
lda.fit(X_train, y_train)

# Extract coeficientes
```

```

coefficients = pd.DataFrame({
    'Feature': X_train.columns,
    'Coefficient': lda.coef_[0],
    'Abs_Coefficient': np.abs(lda.coef_[0])
}).sort_values(by='Abs_Coefficient', ascending=False)

coefficients

```

```
[ ]:           Feature  Coefficient  Abs_Coefficient
12   payment_history_score -53.807819      53.807819
10     job_stability_score -44.244810      44.244810
11   credit_utilization   40.412427      40.412427
2       credit_score      -22.301436      22.301436
4   debt_to_income_ratio   19.025892      19.025892
5       savings_ratio     -14.100416      14.100416
1   employment_years      -11.452017      11.452017
9   residential_stability -8.690222      8.690222
0       annual_income     -7.820056      7.820056
6       asset_value       -6.702356      6.702356
3   open_credit_lines     -6.581518      6.581518
8   education_level      -3.315865      3.315865
7         age             -1.423588      1.423588
15  marital_status_Single  1.099684      1.099684
14  marital_status_Married -0.910776      0.910776
16  marital_status_Widowed  0.702136      0.702136
13  marital_status_Divorced -0.163043      0.163043
```

```
[ ]: lda.classes_
```

```
[ ]: array([0, 1])
```

```
[ ]: lda_scores = lda.transform(X_train)

lda_scores_df = pd.DataFrame(
    lda_scores,
    columns=[f"LD{i + 1}" for i in range(lda_scores.shape[1])],
)

lda_scores_df
```

```
[ ]:          LD1
0   -1.471836
1   -4.005368
2   -1.634797
3   -2.193926
4   -2.033383
...   ...
```

```

1995 -2.590404
1996 -3.692438
1997  8.446930
1998 -2.271978
1999 -3.304125


```

[2000 rows x 1 columns]

```

[ ]: # Show a simple bar chart of your top 3-5 LDA coefficients.
plt.figure(figsize=(10, 6))
plt.bar(coefficients['Feature'].head(5), coefficients['Coefficient'].head(5))
plt.xlabel('Categoría')
plt.ylabel('Valor del coeficiente')
plt.title('Top 5 LDA Coeficientes')
plt.tight_layout()

```



Observando los resultados anteriores, se puede identificar que las variables más importantes para predecir el riesgo default son:

- payment_history_score: -53.886626
- job_stability_score: -44.344218
- credit_utilization: 40.498805
- credit_score: -22.401797
- debt_to_income_ratio: 19.031001

Asimismo, se encontró que cuando el valor del coeficiente es positivo, es significa que cuando la variables incrementa, entonces el riesgo de ser default también. Por el contrario, si este es negativo,

se tiene menos riesgo.

Por lo tanto, las variables que al crecer hacen que haya menos riesgo de tener un estado default, son payment_history_score, job_stability_score y credit_score. Por su parte, las variables que hacen que haya más riesgo son credit_utilization y debt_to_income_ratio.

Esto sugiere que al tener un mejor historial crediticio, tener un empleo estable y un buen score crediticio, la persona tendrá menos riesgo a tener deudas de prestamos. Mientras que si hay un mal uso del credito, esto es si se gasta más de lo que se tiene en el crédito y si se tiene una alta deuda sobre lo que genera, entonces, es hay un riesgo más alto de ser default.

A continuación se realizara el modelo QDA para evaluar si la hipótesis de acepta o se rechaza.

8 Section 6: Model 2 - Quadratic Discriminant Analysis (QDA)

```
[ ]: from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis  
  
# Fit Quadratic Discriminant Analysis  
qda = QuadraticDiscriminantAnalysis()  
qda.fit(X_train, y_train)  
  
/usr/local/lib/python3.12/dist-packages/sklearn/discriminant_analysis.py:1024:  
LinAlgWarning: The covariance matrix of class 0 is not full rank. Increasing the  
value of parameter `reg_param` might help reducing the collinearity.  
    warnings.warn(  
/usr/local/lib/python3.12/dist-packages/sklearn/discriminant_analysis.py:1024:  
LinAlgWarning: The covariance matrix of class 1 is not full rank. Increasing the  
value of parameter `reg_param` might help reducing the collinearity.  
    warnings.warn(  
  
[ ]: QuadraticDiscriminantAnalysis()
```

QDA no produce coeficientes lineales simples, por lo que no es posible realizar una interpretación en este caso.

A continuación se hara una comparación entre ambos modelos, para encontrar el mejor en este caso.

9 Section 7: Model Evaluation & Comparison

```
[ ]: # Se debe de escalar X_test para que quede en las mismas dimensiones  
X_train_fixed = X_train.copy()  
X_train_fixed[columnas_a_escalar] = X_scaled_train  
  
# Escalar columnas númericas  
X_test_scaled_part = pd.DataFrame(  
    scaler.transform(X_test[columnas_a_escalar]),
```

```

        columns=columnas_a_escalar,
        index=X_test.index
    )

# Mantener el resto de las columnas igual
other_cols = [c for c in X_train_fixed.columns if c not in columnas_a_escalar]

# Juntar ambas partes, la escalada y el resto
X_test_pred = pd.concat([X_test_scaled_part, X_test[other_cols]], axis=1)

# Reordenar columnas como en X_train

feature_order = X_train_fixed.columns

X_test_pred = X_test_pred[feature_order]

```

```
[ ]: from sklearn.metrics import ConfusionMatrixDisplay, classification_report,roc_curve, roc_auc_score
from sklearn.metrics import precision_score, recall_score, f1_score
import time
```

```

# Predicciones
start_lda = time.time()
y_pred_lda = lda.predict(X_test_pred)
end_lda = time.time()

start_qda = time.time()
y_pred_qda = qda.predict(X_test_pred)
end_qda = time.time()

# Tiempo de ejecución
lda_time = end_lda - start_lda
qda_time = end_qda - start_qda

print(f"Tiempo de predicción LDA: {lda_time:.6f} segundos")
print(f"Tiempo de predicción QDA: {qda_time:.6f} segundos")

```

```

# Generar matrices de correlación
fig, ax = plt.subplots(1, 2, figsize=(12, 5))
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_lda, ax=ax[0],cmap="Blues")
ax[0].set_title("Confusion Matrix - LDA")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_qda, ax=ax[1],cmap="Oranges")

```

```

ax[1].set_title("Confusion Matrix - QDA")
plt.tight_layout(); plt.show()

# Reporte de Clasificación
print("== Classification Report: LDA ==")
print(classification_report(y_test, y_pred_lda))

print("\n== Classification Report: QDA ==")
print(classification_report(y_test, y_pred_qda))

# Generate the RocCurveDisplay
y_prob_lda = lda.predict_proba(X_test_pred)[:, 1]
y_prob_qda = qda.predict_proba(X_test_pred)[:, 1]

fpr_lda, tpr_lda, _ = roc_curve(y_test, y_prob_lda)
fpr_qda, tpr_qda, _ = roc_curve(y_test, y_prob_qda)
auc_lda = roc_auc_score(y_test, y_prob_lda)
auc_qda = roc_auc_score(y_test, y_prob_qda)


plt.figure(figsize=(7, 6))
plt.plot(fpr_lda, tpr_lda, label=f'LDA (AUC = {auc_lda:.3f})')
plt.plot(fpr_qda, tpr_qda, label=f'QDA (AUC = {auc_qda:.3f})', linestyle="--")
plt.plot([0,1],[0,1], 'k--', lw=1)
plt.xlabel("False Positive Rate"); plt.ylabel("True Positive Rate")
plt.title("ROC Curves: LDA vs QDA"); plt.legend(); plt.grid(True); plt.show()

print(f"AUC (LDA): {auc_lda:.3f}")
print(f"AUC (QDA): {auc_qda:.3f}")

```

Tiempo de predicción LDA: 0.002320 segundos

Tiempo de predicción QDA: 0.002247 segundos




==== Classification Report: LDA ===

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 367 |
| 1 | 1.00 | 1.00 | 1.00 | 133 |
| accuracy | | | 1.00 | 500 |
| macro avg | 1.00 | 1.00 | 1.00 | 500 |
| weighted avg | 1.00 | 1.00 | 1.00 | 500 |

==== Classification Report: QDA ===

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 367 |
| 1 | 1.00 | 1.00 | 1.00 | 133 |
| accuracy | | | 1.00 | 500 |
| macro avg | 1.00 | 1.00 | 1.00 | 500 |
| weighted avg | 1.00 | 1.00 | 1.00 | 500 |



AUC (LDA) : 1.000

AUC (QDA) : 1.000

10 Section 8: Technical Conclusion & Model Selection

Observando, los resultados, se observa que ambos modelos logran clasificar el estatus de deuda al 100%, en todas las métricas. Asimismo, se midió el tiempo de ejecución de ambos modelos para observar cuál es más rápido. En este caso, se puede observar que el modelo que tarda menos es el de QDA, aunque la diferencia de tiempo es pequeña, por lo que no se puede concluir que el modelo es significativamente más rápido.

Por lo tanto, se rechaza la hipótesis, ya que a pesar de tener una matriz de covarianza diferente, el modelo LDA se comportó de la misma manera que el QDA. Por lo tanto, la elección de un modelo es indiferente; ambos resultados dan lo mismo. Sin embargo, en caso de que la empresa busque eficiencia, el mejor modelo es QDA, aunque el tiempo de diferencia es relativamente el mismo.

Es importante mencionar que en este caso lo que se quiere reducir el número de falsos negativos, ya

que se busca evitar que los clientes con deudas, sean seleccionados como personas sin deudas. De igual manera, se busca reducir el número de falsos positivos, ya que esto haría que la empresa no seleccionara a clientes sin deuda. Lo que sería un resultado no deseado, ya que la empresa busca reducir la probabilidad de que las personas tengan una deuda default. Por lo tanto, la métrica que tiene más importancia en este caso es F1-score, ya que se busca tener un balance harmónico entre la precisión y sensibilidad.

Por ende, en este caso, es de mayor importancia observar el f1_score, aunque como se mencionó anteriormente, en este caso ambos modelos lograron un resultado del 100%.

...